

## **Kubernetes Processes**

There is a core set of *Kubernetes* processes that form the cluster runtime. The cluster runtime is further divided into master components, aka the control plane, and node components. There is also a set of cluster add-ons.

### **Master Components**

#### **kube-apiserver**

The kube-apiserver exposes the *Kubernetes* API. It is where the *Kubernetes Resources* are published into.

#### **kube-controller-manager**

The kube-controller-manager runs a number of controllers that are responsible for watching for particular *Kubernetes Resources* in the API. The controllers bring the cluster into alignment with the declared

state in the *Resources*.

## **kube-scheduler**

The kube-scheduler assigns *Pods* to nodes based on a number of criteria. Such as resource requests and resource availability, workload placement policies, etc.

## **cloud-controller-manager**

The cloud-controller-manager runs controllers that interact with a specific cloud provider. An example would be *AWS* or *Digital Ocean*.

## **Etcd**

*Etcd* is a distributed key / value store. It is used to store the *Kubernetes* cluster state.

## **Node Components**

### **kubelet**

The kubelet runs on each node and is responsible for managing the *Pods* on that node.

### **kube-proxy**

The kube-proxy configures the network

rules so that abstractions like the *Service Resource* work. *Services* are described in some detail further on in the document.

## Kubernetes Concepts

Before describing *Kubernetes Resources* there is some important conceptual information to cover.

### Labels and Selectors

Labels is one of the fundamental *Kubernetes* concepts. A label is a key / value pair that is attached to a *Kubernetes Resource*. And each *Kubernetes Resource* can have many labels. Labels serve as a means to identify a particular *Resource* or set of *Resources*.

Selectors are also key / value pairs. They are used to select *Resources* that have a particular label or set of labels. By this mechanism it is possible for one *Resource* type to select one or more *Resources* of another type. Labels can also be used with *kubectl* to query just particular *Resources*.

The following diagram depicts this selection mechanism with a *Service* selecting a set of *Pods*.

More information about labels and selectors can be found [here](#).

### Kubernetes Resources

*Resources* are high level abstractions that contain the declarative state for an infrastructure component. *Resources* are published into the *Kubernetes* API server and it is the responsibility of *Kubernetes* controllers to bring the cluster into agreement with the declared state.

There are many different types of *Resource* available. The full list of *Resources* can be found in the [API documentation](#). The following sections describe just a few of the most important *Resources*.

#### Note

This is really just a high level introduction. More detailed information about some of these *Resources* will be provided as the course unfolds.

### Workloads

#### Pods

As described previously, a *Pod* is the fundamental unit of execution in *Kubernetes*. A *Pod* has one or more containers. The following diagram depicts a *Pod* and the three types of containers that it can have.

Container Type	Description
----------------	-------------

Application	This is the core container for an application. You must have one of
-------------	---

these with the most common pattern being that a *Pod* only has the application container.

**Sidecar** A *Pod* can also have other containers that perform some useful work in support of the application container. A container that ships application logs is a good example. There can be any number of these within a *Pod*.

**Init** Sometimes it is necessary to do some initialisation before starting the application container. For instance, to create a database. There can be any number of these *Init*containers within a *Pod* and they are run one at a time in sequential order. See [here](#) for more information.

As with other *Resource* types, a *Pod* can have labels. And these labels are used by other *Resources* (see next section) to enable management of the *Pods*.

## Deployments

A *Deployment* allows you to manage the lifecycle of *Pods* and an associated *Resource* called a *ReplicaSet*.

A *Deployment* contains a specification for a *Pod* and also additional information, such as the number of *Pods* to run.

The *ReplicaSet* is created when a *Deployment* is created or updated and it is actually the *ReplicaSet* that is used as the definition for creating the *Pods*.

Each time the *Deployment* is updated a new *ReplicaSet* is created. This makes it possible to roll back a *Deployment* by using the previous *ReplicaSet*.

The following diagram depicts this.

More informations about *Deployments* can be found [here](#).

## DaemonSet

A *Daemonset* ensures that a *Pod* is run on all, or a subset, of the available *Kubernetes Nodes*.

A use case for this might be where you want to have a metric collector run on all the *Nodes* so that kernel performance statistics can be gathered.

The following diagram depicts this.

More information about *DaemonSets* can be found [here](#).

## StatefulSet

A *StatefulSet* manages sets of *Pods* in a way that provides a guaranteed *Pod* name and a specific storage volume for each. The order in which the *Pods* are created and destroyed is also guaranteed.

A use case for this is where a distributed database cluster relies upon its members having a specific name with associated data and that the cluster must be created with an initial bootstrap member. *ElasticSearch* is a good example of this.

The following diagram depicts this.

## Job

A *Job* creates one or more *Pods* and ensures that at least a certain amount of them run to completion.

A use case for this is a batch data load.

The following diagram depicts this.

## **CronJob**

A *CronJob* runs *Jobs* on a schedule. The format for the schedule is the **Cron** format.

A use case for this is a scheduled backup.

The following diagram depicts this.

More information about *CronJobs* can be found [here](#).

## **Discovery and Load Balancing**

### **Service**

A *Service* is used to group together one or more *Pods* using a selector.

The *Service* then provides a *ClusterIP* and a DNS name for the *Pods*.

The *Kube Proxy* process uses the information contained within a *Service* to configure appropriate iptables rules on each of the nodes so that clients of the *Service* are able to be routed to an appropriate *Pod*.

There is a lot more detail to *Services* which can be found in the link below the diagram.

## **Ingress**

An *Ingress* is used to configure an *Ingress Controller* so that a *Service* can be made accessible externally of the cluster.

An *Ingress* is a layer 7 HTTP construct and can configure both HTTP and HTTPS access.

An *Ingress Controller* is a reverse proxy that watches for *Ingress Resources* in the *Kubernetes* API server and configures itself with appropriate rules as per the *Ingress*.

## **Config and Storage**

### **ConfigMap**

A *ConfigMap* allows configuration files to be stored independently of a container image. The files stored within the *ConfigMap* can then be mounted into a container within a *Pod* at runtime.

The following diagram depicts this.

### **Secret**

A *Secret* allows secret files and values to be stored independent of a container image. Much like a *ConfigMap* but also with the ability to store the value of a key in an environment variable.

The following diagram depicts this.

More information about *Secrets* can be found [here](#).

### **PersistentVolume and PersistentVolumeClaim**

A *PersistentVolume* represents some form of distributed storage.

A *PersistentVolumeClaim* allows a *PersistentVolume* to be claimed for a *Pod*.

The following diagram depicts this.

### **Namespace**

This is where we put it all together. A *Namespace* can be thought of as like an environment. Collections of *Resources* can be deployed into the *Namespace* to form a grouping of related components.

Examples of *Namespace* use might be to create a DEV environment, or to share a set of services that are used by other *Namespaces*, i.e. an *Ingress Controller*.

The following diagram depicts this.